

Benchmark Report: nbody

Five bodies, twenty thousand steps, and a sixth of the runtime spent computing the number zero

Matan Cohen · Yuval Kogan · HWSW Final Project, Technion · github.com/matanco64/HWSW_Project

A probe is falling toward Neptune, and somebody has to know where it will be in six hours. That is this benchmark: integrate the Sun and the four gas giants forward under Newtonian gravity, twenty thousand steps at a time. It is thirty lines of arithmetic with no libraries, no I/O, and no data structure more exotic than a list of floats — which makes it the cleanest possible place to ask a question with a surprising answer: **when CPython runs this, what is it actually doing?** Almost none of it is gravity. Roughly a sixth of the runtime goes into converting the constant 0 into a C array index, twenty thousand times over, and that observation is what the optimization in §3 and the accelerator in §5 are both built on.

1. Overview

nbody comes from the Computer Language Benchmarks Game: a **symplectic-Euler integration** of the Sun plus Jupiter, Saturn, Uranus and Neptune under Newtonian gravity. Five bodies means **ten interacting pairs**, computed once at import and reused for the whole run. Each pyperf iteration runs `report_energy()`, `advance(dt=0.01, n=20000)`, then `report_energy()` again.

- **Libraries:** `stdlib` only — the stock version does not even import `math`. Nothing is vectorized, nothing is compiled, there is no I/O in the timed region.
- **Data structures:** a list of `[position, velocity, mass]` triples, where position and velocity are 3-element lists of floats; `PAIRS` is a list of 2-tuples of those triples. Total mutable state is **35 floats**.
- **Correctness oracle:** total system energy. A symplectic integrator conserves it to $O(dt)$, so any change to the physics shows up immediately.




Figure 1: The entire workload. Because $N = 5$, the pair list is **fixed for the whole run** — the fact §3 exploits.

2. Initial analysis

Measurement setup. Course QEMU VM: Ubuntu 22.04, CPython 3.10.12, pyperformance 1.14.0, KVM on a Technion host. Timings are `pyperformance run - - rigorous` (40 processes / 120 values) on the release `python3`. Profiles are recorded separately with `python3-dbg` for symbols and are **never** quoted as timings — the debug build’s assertions and allocator hooks inflate interpreter-internal frames roughly 2.5–3×. Two guest quirks worth recording: the `cycles` PMU event records **zero** samples under KVM, so sampling must use `-e cpu-clock`; and `perf stat` silently returns zeros past four events, so counters are taken in two passes.

Stock **nbody** measures **231 ms ± 3 ms**, reproduced to within 1 ms across runs a week apart — an unusually stable benchmark. The flame graph is a single `advance()` plateau over `_PyEval_EvalFrameDefault`, about 95% of samples. The interesting part is underneath it:

Cost group	Self	Symbols
List subscript machinery	~15%	<code>list_ass_item</code> 2.83, <code>PyNumber_AsSsize_t</code> 2.12, <code>_PyNumber_Index</code> 1.87, <code>PyObject_GetItem</code> 1.65, <code>list_item</code> 1.65, <code>PyObject_SetItem</code> 1.47, <code>PyLong_AsSsize_t</code> 1.43, <code>list_ass_subscript</code> 1.33
Generic binary-op dispatch	~8%	<code>binary_op1</code> 5.75, <code>_Py_CheckSlotResult</code> 2.22
Float boxing	large	<code>PyFloat_FromDouble</code> / <code>float_dealloc</code> churn
Actual arithmetic	small	<code>libm pow</code> for <code>d**-1.5</code> , plus the adds and multiplies

On CPython 3.10 every `v1[0] -= dx * b2m` is a full `PyObject_GetItem` → `_PyNumber_Index` → `PyLong_AsSsize_t` → `list_item` round trip, and the reverse for the store. **None of that is physics.** CPython 3.11 later added `BINARY_SUBSCR_LIST_INT` specialization to make exactly this cheap; 3.10 — the version the course VM runs — has none of it.

ph --nonblocking -o /home/ubuntu/hwsproject/results/pyspy_nbody_stock.svg -- python3 /usr/local/lib/python3.10/dist-packages/pyperformance/data-files/benchmarks/t

<module> (run_benchmark.py:155)
bench_time_func (pyperf/_runner.py:502)
_main (pyperf/_runner.py:463)
_worker (pyperf/_runner.py:437)
create_run (pyperf/_worker.py:300)
compute (pyperf/_worker.py:364)
compute (pyperf/_worker.py:289)
compute_warmups_values (pyperf/_worker.py:271)
_compute_values (pyperf/_worker.py:79)
task_func (pyperf/_runner.py:497)
bench_nbody (run_benchmark.py:132)
advance.. adv.. a.. a.. advance (run_benchmark.py.. ad.. advance (r.. advance (r.. advance (r.. advance (run_b.. advan.. advance (run_b.. advan.. advance.. ad.. advan..

Figure 2: **Before, Python frames** (py-spy). CPython 3.10 has no `-X perf` trampoline, so `perf` cannot name Python functions; py-spy samples the interpreter's own frame stack. `advance` is the single wide plateau under `bench_nbody`.

A note on how these were recorded. Our first captures used `perf record -g` (frame-pointer unwinding) and were **unusable**: 1,877 [unknown] frames and call chains full of return addresses like `0xfdfdfdfdfdfdf00` — the `Py_DEBUG` freed-memory fill byte. Ubuntu’s `python3-dbg` is built without frame pointers, so the unwinder was walking into freed stack memory. `--call-graph dwarf,16384` fixes it (0 such addresses, 105 unknowns) at the cost of 100× larger captures and 20 minutes per profile. The flat self-percentages above were always valid; only the call graph was affected.

3. Optimizations

3a. What shipped: compile the schedule away

The inner loop rediscovers, twenty thousand times, a schedule that is **constant for the entire run**: which bodies interact, in what order, with which masses. Only the 30 coordinate floats change. So we compute that schedule once, at import time and outside the timed region, and emit straight-line source for `advance(dt, n)` in which every coordinate is a Python **local** (`LOAD_FAST`) and every mass is a **literal** (`LOAD_CONST`). State is read out of the body lists before the step loop and written back after, so `report_energy()` observes the same objects. This is **partial evaluation** — the technique `dataclasses` and Django’s template compiler use.

Per integration step	Stock	Generated
Bytecodes executed	1484.5	850.6
List subscripts (<code>BINARY_SUBSCR/STORE_SUBSCR</code>)	150	4

Removing 146 of 150 subscripts per step attacks precisely the ~15% tower in §2. The bytecode counts come from `sys.monitoring` and are a **version-independent structural metric**: unlike wall clock, they do not depend on which CPython specializes what.

The output is bit-for-bit identical. All 35 state floats compare `==` against stock after 20,000 steps, and so does `report_energy()`: `max |Δ| = 0.0` — an equality, not a tolerance. We kept the stock `dt * (dsq ** -1.5)` rather than the cheaper `dt / (dsq * sqrt(dsq))` to preserve that claim, and measurement justified the choice independently: the `sqrt` variant was not distinguishable from zero on any interpreter tested, and on one run came out slower. The emitter is **generic in N** — hand it nine bodies and it emits a nine-body kernel — so this specializes the given workload rather than hard-coding an answer.

3b. Three anti-results

- **Fast Multipole / Barnes–Hut is a 3.4× slowdown at N = 5.** We built a real octree Barnes–Hut (monopole, $\theta = 0.5$, tree rebuilt every step, validated by driving $\theta \rightarrow 0$ to reproduce direct summation to $1.2e-16$) and measured the crossover: BH/direct is **3.39** at $N = 5$, 1.58 at $N = 100$, and only drops below 1.0 near $N \approx 300\text{--}400$ — this benchmark runs 70× below that. At $N = 5$ the tree performs 18 force-term evaluations where direct summation performs 10, and builds a 7-node tree 20,000 times on top. Note also that $\log_2 5 = 2.3$, so “ $O(N \log N)$ beats $O(N^2)$ ” compares $5 \times 2.3 \approx 12$ against $25/2 \approx 12$: **there is no asymptotic gap to exploit at this size**. Barnes–Hut is the *cheaper* tree code, so this is a strict lower bound on FMM’s constant factor. `pyperformance` itself ships `bm_barnes_hut` as a *separate* benchmark — upstream treats the tree code as a different workload, not an optimization of this one.
- **Struct-of-arrays loses (0.88×).** Textbook advice, wrong here. In C, `x[i]` is a pointer add and the lanes vectorize; in CPython it is a `PyObject_GetItem` round trip, and the stock layout already hands the loop two *direct* velocity-list locals. SoA raises subscripts per step from 150 to 245. A data-layout optimization is only an optimization relative to a cost model, and the interpreter’s cost model is not the hardware’s.
- **NumPy loses (0.35×).** At vector length 5–25 there is nothing to amortize the per-call dispatch against.

We also disproved a widely-repeated micro-optimization recorded in our own research notes: “stop destructuring positions, index `r1[0] - r2[0]` directly”. Measured against a semantic-no-op control, it is worth **1.005× — nothing**. The real win in that tier (1.12×) came from a different change: turning six subscript read-modify-writes into six plain stores.

4. Performance comparison

Stock and optimized were run back-to-back in the same session on the VM with `pyperformance run --rigorous`, the optimized build through a custom `--manifest` that keeps the benchmark name identical so `pyperf compare_to` matches by name. Significance is Student’s two-tailed t-test at 95% confidence.

Configuration	Mean ± std dev	Speedup	Correctness
Stock <code>pyperformance</code> 1.14.0	231 ms ± 3 ms	1.00×	reference
Optimized (shipped, pure Python)	141 ms ± 4 ms	1.64×	bit-identical state + energy

A **39.0% reduction in runtime** against a requirement of 7% — **5.6× the bar** — and the difference is reported significant. On CPython 3.12 the same change measures ~1.5× rather than 1.64×: 3.11+ adaptive specialization already recovers part of the subscript overhead we remove, so the win is *larger* on the 3.10 the course targets. Re-running on a newer interpreter should be expected to give a smaller number, and that is not a regression.

energy oracle — an experiment our bit-identical software baseline makes meaningful. *rsqrt*: ROM seed width against Newton–Raphson iteration count trades table area for latency; this is the operation that still dominates even the compiled Rust kernel, so it is where fixed-function hardware earns its win over software. On power, the GRAPE lineage is the evidence that a fixed-function gravity pipeline reaches roughly an order of magnitude better Gflops/W than a general-purpose core on this kernel.

6. Conclusion

The benchmark looks like a physics problem and behaves like an interpreter problem. About a sixth of stock runtime goes into turning the constant `0` into an array index, which is why the change that paid was not a better algorithm but the removal of 146 of 150 list subscripts per step: **1.64× on the course VM, 39% less runtime against a 7% requirement, with bit-for-bit identical output.** The Rust kernel takes the same idea to its conclusion at 24.1×, still bit-identical, and in doing so exposes the one remaining arithmetic bottleneck that only hardware fixes.

The negative results were worth more than the win. Barnes–Hut — and by extension FMM — is **3.4× slower** at $N = 5$ with a crossover near $N \approx 350$; struct-of-arrays, the standard data-layout advice, *loses* in an interpreter; NumPy loses at this vector length; and a widely-repeated micro-optimization is worth statistically nothing. Each cost less to establish than the effort it redirected. The same discipline paid a second time in hardware: the measured 95% share of `advance()` is what makes a step-count register worth building, and the measured residual cost of `pow` is what puts the *rsqrt* unit at the centre of the design.